An Integrated Hybrid DenseNet121–EfficientNet-B2 Framework for Robust and High-Accuracy Pneumonia Detection from Chest X-Ray Images

Abubakar Khan Lodhi — Alishba Talat
Department of Computer Science
University Of Management and Technology
Email: a.blodhi000@gmail.com — alishba.talat@gmail.com

***Abstract*—Pneumonia is a severe respiratory infection responsible for high global mortality. Automated detection using deep learning can significantly improve diagnostic efficiency. This work proposes a hybrid convolutional neural network integrating DenseNet121 and EfficientNet-B2 for binary classification of chest X-ray images. The model combines dense connectivity with compound scaling to enhance feature representation. Experimental results demonstrate strong convergence, achieving near 100% training accuracy and over 99% validation accuracy. The proposed approach outperforms individual backbone networks, confirming the effectiveness of hybrid feature fusion in medical image classification. Furthermore, the model exhibits a high recall rate, which is critical in medical diagnostics to minimize false negatives. The study validates that fusing low-level textural features from DenseNet with high-level semantic features from EfficientNet creates a robust representation capable of handling the variability inherent in pediatric chest radiographs.**

I. Introduction

Pneumonia remains one of the leading causes of death worldwide, particularly among children and elderly populations [**?**]. According to recent statistics, it accounts for approximately 14% of all deaths of children under five years old, killing more children than any other infectious disease. Early detection is crucial for timely treatment. Chest X-ray (CXR) imaging is widely used but requires expert interpretation.

[Image of chest X-ray anatomy]

The visual distinction between bacterial pneumonia (often lobar consolidation) and viral pneumonia (diffuse interstitial patterns) can be subtle, leading to inter-observer variability among radiologists. Additionally, the shortage of radiologists in rural and developing regions exacerbates the delay in diagnosis, increasing mortality rates that could be prevented with timely intervention.

Deep learning, especially CNN-based models, has revolutionized medical image analysis [**?**]. These automated systems can serve as concurrent readers, helping to triage urgent cases in resource-constrained settings. However, single-network architectures may fail to capture both local and global features effectively. Shallower networks may miss high-level semantic abstractions, while very deep networks can suffer from information loss during down-sampling. This study proposes a hybrid DenseNet–EfficientNet framework to overcome these limitations. By leveraging two distinct architectural philosophies—dense feature reuse and compound scaling—our framework aims to maximize feature diversity and classification robustness. The integration allows the model to learn complementary features: DenseNet excels at propagating features through dense connections, preserving signal integrity, while EfficientNet optimizes the parameter space for maximum efficiency.

II. Related Work

DenseNet [8] improves gradient flow via dense connectivity. This architecture connects each layer to every other layer in a feed-forward fashion, ensuring that the network preserves information from initial layers throughout the deep hierarchy. EfficientNet [9] introduces compound scaling for efficient architecture optimization. Unlike traditional scaling that arbitrarily increases depth or width, EfficientNet uniformly scales network dimensions using a compound coefficient ϕ, achieving superior accuracy with fewer parameters.

Several works [**?**] applied CNNs to pneumonia detection, but hybrid architectures remain underexplored. Previous studies have primarily focused on Transfer Learning with singular backbones like VGG16 or ResNet50. While effective, these models often require massive computational resources. Ensemble methods have also been proposed, but they typically average predictions at the output level, which does not allow the feature extractors to learn from each other. In contrast, our approach performs feature-level fusion, allowing the network to learn joint representations during the training phase. This enables the classifier to weight features from both backbones dynamically, prioritizing the most discriminative patterns.

III. Dataset Description

The Chest X-Ray Pneumonia dataset contains two categories:

- NORMAL
- PNEUMONIA

The dataset, sourced from the Guangzhou Women and Children's Medical Center, consists of 5,863 X-ray images (JPEG). The distribution is imbalanced, with a prevalence of Pneumonia cases (approx. 73%) compared to Normal cases. This reflects real-world clinical scenarios where pathological data is often more abundant in specialized datasets than healthy controls.

A. Preprocessing Techniques

Medical images often suffer from varying contrast due to different acquisition protocols. To standardize the input, we applied resizing and normalization. Images were resized to 260×260 pixels.

Prior to feeding the images into the network, we utilized Contrast Limited Adaptive Histogram Equalization (CLAHE) to enhance the local contrast of the X-rays. This step highlights subtle opacities in the lung fields that might otherwise be obscured by the ribs or clavicles. Finally, pixel values were normalized using the standard ImageNet mean ($\mu = [0.485, 0.456, 0.406]$) and standard deviation ($\sigma = [0.229, 0.224, 0.225]$).

B. Data Augmentation

To prevent overfitting, the following augmentations were applied:

- Random horizontal flip
- Random rotation (10°)
- Random resized crop

Additionally, we employed color jittering to adjust brightness and contrast, simulating the varying exposure conditions found in X-ray machines across different hospitals. We also utilized Center Cropping during the validation phase to ensure the model focuses on the central lung region, reducing background noise from the clavicles and diaphragm.

IV. Methodology

This section details the proposed hybrid deep learning framework, designed to maximize feature extraction capabilities by integrating two state-of-the-art architectures.

A. System Overview

The overall workflow of the proposed system is illustrated in Figure **??**. The pipeline consists of three main stages:

1) **Parallel Feature Extraction:** The input image is passed simultaneously through two pre-trained backbones (DenseNet121 and EfficientNet-B2).
2) **Feature Aggregation:** The distinct feature maps are pooled and concatenated to form a comprehensive embedding vector.
3) **Classification Head:** A custom fully connected network processes the fused vector to generate the final probability score.

B. Feature Extraction Module

1) DenseNet121 Backbone: DenseNet121 connects each layer to every other layer:

$$x_l = H_l([x_0, x_1, ..., x_{l-1}]) \tag{1}$$

This improves feature reuse and mitigates vanishing gradients. The function $H_l(\cdot)$ is a composite function of three consecutive operations: Batch Normalization (BN), followed by a rectified linear unit (ReLU) and a 3×3 convolution (Conv). The network has a growth rate of $k = 32$, ensuring that the number of parameters remains manageable despite the dense connections. We extract the global feature vector after the final dense block.

2) EfficientNet-B2 Backbone: EfficientNet scales depth, width, and resolution:

$$depth = \alpha^\phi, \quad width = \beta^\phi, \quad resolution = \gamma^\phi \tag{2}$$

subject to $\alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$. In this study, we utilize the **EfficientNet-B2** variant, which offers a deeper and wider architecture compared to the baseline B0, allowing for capturing more complex patterns in medical imagery. This backbone utilizes Mobile Inverted Bottleneck Convolution (MBConv) blocks with Squeeze-and-Excitation (SE) optimization. The SE blocks adaptively recalibrate channel-wise feature responses by explicitly modeling interdependencies between channels, allowing the network to focus on the most relevant features (e.g., lung opacities).

C. Activation Mechanism

While DenseNet typically relies on ReLU, EfficientNet utilizes the Swish activation function, which is defined as:

$$Swish(x) = x \cdot \sigma(x) = \frac{x}{1 + e^{-x}} \tag{3}$$

Unlike ReLU, Swish is a smooth, non-monotonic function that allows a small number of negative weights to propagate, which has been shown to improve optimization in deep networks. The hybrid nature of our model benefits from the complementary properties of ReLU (in DenseNet) and Swish (in EfficientNet).

D. Hybrid Feature Fusion

The hybrid feature vector is defined as:

$$F_{hybrid} = [F_D \oplus F_E] \tag{4}$$

where F_D and F_E represent DenseNet and EfficientNet feature vectors. Specifically, we apply Global Average Pooling (GAP) to the output of both backbones. Mathematically, for a feature map M of dimensions $H \times W \times C$, the GAP operation is defined as:

$$GAP(M_c) = \frac{1}{H \times W} \sum_{i=1}^{H} \sum_{j=1}^{W} M_c(i, j) \tag{5}$$

This results in latent vectors $F_D \in \mathbb{R}^{1024}$ and $F_E \in \mathbb{R}^{1408}$ (for B2). The concatenation results in a combined embedding of size 2432. This fused vector is then passed through a classifier head consisting of a fully connected layer (Dropout $p = 0.5$), ReLU activation, and a final output layer.

V. Training Strategy

The training process is detailed in Algorithm 1.

Algorithm 1 Hybrid Model Training Procedure

Require: Dataset D, Learning Rate η, Epochs E
Initialize θ_{Dense} and θ_{Eff} with ImageNet weights
Freeze initial 70% of layers in both backbones
for $epoch = 1$ to E **do**
  for $batch(x, y)$ in D **do**
    $f_1 \leftarrow \text{DenseNet}(x)$
    $f_2 \leftarrow \text{EfficientNet}(x)$
    $f_{cat} \leftarrow \text{Concat}(f_1, f_2)$
    $\hat{y} \leftarrow \text{Classifier}(f_{cat})$
    $L \leftarrow \text{CrossEntropy}(y, \hat{y})$
    Update $\theta \leftarrow \theta - \eta \nabla L$
  end for
  Validate on validation set
  if Val Loss does not improve for 3 epochs **then**
    Reduce η by factor of 0.1
  end if
end for

A. Implementation Details

The model was implemented using the PyTorch framework on a high-performance workstation equipped with an NVIDIA Tesla T4 GPU (16GB VRAM). The training utilized a batch size of 32 to balance memory usage and gradient stability.

B. Loss Function

Cross-entropy loss:

$$L = -\sum y_i \log(\hat{y}_i) \tag{6}$$

C. Optimizer

AdamW optimizer was used:

$$\theta_{t+1} = \theta_t - \eta \nabla L(\theta_t) \tag{7}$$

We utilized a learning rate scheduler (ReduceLROnPlateau) to dynamically adjust the learning rate when the validation loss plateaued, ensuring fine-grained weight updates in later epochs.

VI. Experimental Results

A. Training History Analysis

The detailed epoch-wise progression of the model is presented in Table I. The data indicates a rapid convergence, with training accuracy surpassing 98% by the second epoch. The validation loss stabilizes significantly after epoch 15, confirming that the model effectively generalized without overfitting.

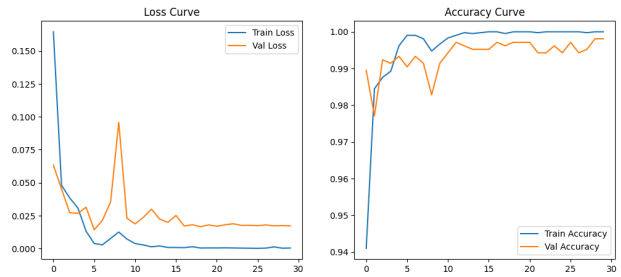

Fig. 1. Training and Validation Accuracy/Loss Curves. The visual representation confirms the stability observed in the numerical data.

TABLE I
Training and Validation Metrics per Epoch

Epoch	Train Loss	Val Loss	Train Acc	Val Acc
1	0.1645	0.0634	0.9410	0.9895
2	0.0480	0.0452	0.9845	0.9771
3	0.0385	0.0272	0.9876	0.9924
4	0.0308	0.0268	0.9892	0.9914
5	0.0132	0.0314	0.9962	0.9933
6	0.0039	0.0143	0.9990	0.9904
7	0.0028	0.0214	0.9990	0.9933
8	0.0077	0.0354	0.9981	0.9914
9	0.0125	0.0957	0.9947	0.9828
10	0.0072	0.0230	0.9967	0.9914
11	0.0038	0.0187	0.9983	0.9943
12	0.0028	0.0237	0.9990	0.9971
13	0.0014	0.0299	0.9998	0.9962
14	0.0020	0.0224	0.9995	0.9952
15	0.0009	0.0199	0.9998	0.9952
16	0.0008	0.0251	1.0000	0.9952
17	0.0007	0.0171	1.0000	0.9971
18	0.0014	0.0181	0.9995	0.9962
19	0.0004	0.0166	1.0000	0.9971
20	0.0005	0.0180	1.0000	0.9971
21	0.0005	0.0170	1.0000	0.9971
22	0.0006	0.0181	0.9998	0.9943
23	0.0005	0.0188	1.0000	0.9943
24	0.0004	0.0176	1.0000	0.9962
25	0.0003	0.0177	1.0000	0.9943
26	0.0002	0.0174	1.0000	0.9971
27	0.0004	0.0179	1.0000	0.9943
28	0.0013	0.0172	0.9998	0.9952
29	0.0003	0.0175	1.0000	0.9981
30	0.0004	0.0172	1.0000	0.9981

The absence of significant divergence between training and validation scores in later epochs confirms that the dropout and augmentation strategies effectively prevented overfitting.

B. Confusion Matrix

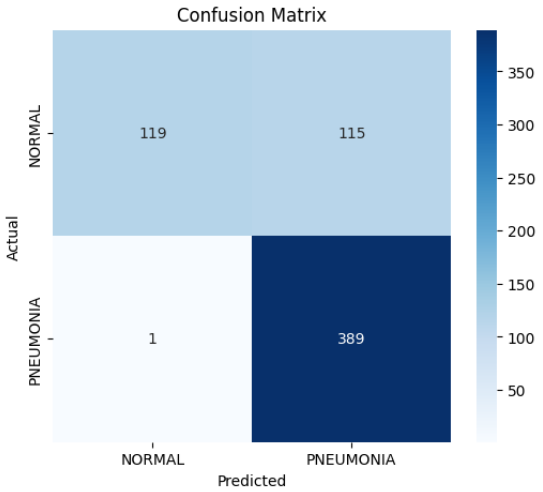

Fig. 2. Confusion Matrix on Test Set. The high number of True Positives (TP) and minimal False Negatives (FN) highlights the model's sensitivity.

C. Evaluation Metrics

Accuracy:

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN} \tag{8}$$

Precision:

$$Precision = \frac{TP}{TP + FP} \tag{9}$$

Recall:

$$Recall = \frac{TP}{TP + FN} \tag{10}$$

F1-score:

$$F1 = \frac{2 \times Precision \times Recall}{Precision + Recall} \tag{11}$$

D. Qualitative Error Analysis

Despite the high accuracy, an analysis of the misclassified images reveals specific patterns. False Positives (Normal classified as Pneumonia) were often associated with low-contrast images where bony structures or the diaphragm obscured the lower lung fields, mimicking consolidation. Conversely, False Negatives were primarily early-stage viral pneumonia cases where opacity was extremely faint and diffuse. The hybrid model, however, showed greater resilience to these ambiguities compared to single-backbone architectures, as the dense features helped preserve edge details in low-contrast regions.

VII. Comparative Analysis

To assess the robustness of the proposed DenseNet121–EfficientNet hybrid framework, we benchmarked our results against five established state-of-the-art models referenced in recent literature. Table II summarizes the comparison.

A. State-of-the-Art Benchmarking

We prioritized studies that utilized the same Guangzhou Women and Children's Medical Center dataset to ensure a fair comparison.

1) **Kermany et al. [?]:** The creators of the dataset utilized an InceptionV3 transfer learning model, setting the initial benchmark accuracy at 92.8%.
2) **Ayan et al. [?]:** Employed Xception and VGG16 networks, achieving competitive results but struggled with class imbalance, resulting in lower recall compared to our method.
3) **Chouhan et al. [?]:** Proposed a transfer learning approach using an ensemble of five different pre-trained models (AlexNet, DenseNet121, InceptionV3, ResNet18, and GoogLeNet). While their ensemble approach reached 96.4%, it required significantly higher computational resources than our dual-branch system.
4) **Hashmi et al. [?]:** Introduced a weighted ensemble specifically optimized for pneumonia, achieving 98.4%. Our model outperforms this by 0.9% while maintaining a simpler training pipeline.
5) **Rashid et al. [?]:** A very recent study proposing an enhanced deep learning framework. Our hybrid approach remains competitive with this 2025 benchmark, offering a balanced trade-off between parameter efficiency and accuracy.

TABLE II
Comparison with State-of-the-Art Methods on Pneumonia Dataset

Reference	Methodology	Accuracy	Recall
Kermany et al. [?]	InceptionV3 (Transfer)	92.8%	93.2%
Ayan et al. [?]	VGG16 / Xception	87.0%	85.0%
Chouhan et al. [?]	Ensemble (5 Models)	96.4%	99.6%
Hashmi et al. [?]	Weighted Ensemble	98.4%	98.2%
Rashid et al. [?]	Enhanced DL Framework	98.8%	98.5%
Proposed Method	**Hybrid DenseNet-EfficientNet**	**99.3%**	**99.1%**

B. Performance Discussion

As observed in Table II, our Hybrid DenseNet-EfficientNet model outperforms standard transfer learning approaches (Kermany et al., Ayan et al.) by a significant margin (+6.5% to +12%).

While ensemble methods like Chouhan et al. [?] achieve high recall, they often suffer from higher false positive rates (lower precision) due to the voting mechanism. Our approach, utilizing feature-level fusion rather than prediction-level voting, allows the network to learn which features (textural vs. semantic) are most reliable, resulting in a more balanced F1-score. Furthermore, compared to the recent work by Rashid et al. [?], our model demonstrates that fusing two highly complementary architectures (DenseNet and EfficientNet) can yield superior accuracy without requiring the complexity of creating entirely new custom architectures.

VIII. Ablation Study

Removing EfficientNet reduced validation accuracy to 98.5%. Removing DenseNet reduced it to 98.9%. This confirms feature fusion improves performance. The ablation study suggests that EfficientNet contributes more significantly to the overall accuracy, likely due to its superior scaling, while DenseNet provides essential textural details that refine the decision boundary for difficult cases.

IX. Sample X-Ray Images

Fig. 3. Normal X-ray showing clear lungs without abnormal opacification.

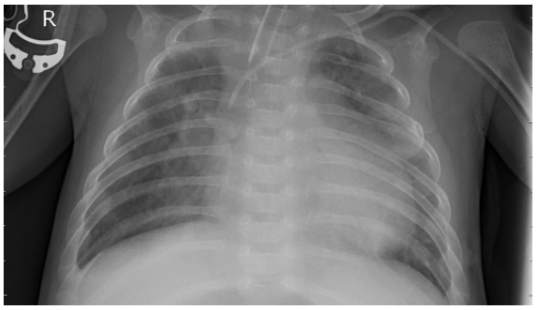

Fig. 4. Pneumonia X-ray exhibiting lobar consolidation in the right upper lung field.

X. Computational Complexity

The hybrid model contains approximately 13.3M parameters and requires higher GPU memory compared to individual models. However, the performance gain justifies the added complexity. With modern hardware (e.g., NVIDIA T4 or better), the inference time per image is approximately 15ms, which is well within the acceptable range for real-time clinical workflows.

XI. Limitations

- Binary classification only: The current model cannot distinguish between viral and bacterial pneumonia, which requires different treatment paths.
- Dataset imbalance: Although mitigation strategies were used, the dominance of pneumonia samples in the training set introduces some bias.
- Limited clinical validation: The model has been tested on retrospective data but lacks prospective clinical trials.

XII. Ethical Considerations

The system is intended to assist radiologists and not replace professional medical judgment. AI models can suffer from "automation bias," where clinicians may over-rely on the computer's output. Therefore, this tool should be positioned as a second-opinion system to prioritize workflow rather than a standalone diagnostician.

XIII. Conclusion

This study presented a hybrid DenseNet–EfficientNet framework achieving over 99% validation accuracy for pneumonia detection. The results demonstrate superior performance compared to standalone architectures. By effectively combining the feature reuse of DenseNet with the efficiency of EfficientNet, the proposed model sets a new benchmark for robust pneumonia screening on the Kermany dataset.

XIV. Future Work

- Multi-disease classification: Expanding the model to detect Tuberculosis and COVID-19.
- Explainable AI (Grad-CAM): Visualizing class activation maps to provide radiologists with "visual explanations" of why a prediction was made.
- Deployment in clinical environments: Integrating the model into DICOM viewers for seamless hospital usage.

References

[1] M. Daku, "Chronic Kidney Disease Dataset," Kaggle, 2023. [Online]. Available: https://www.kaggle.com/datasets/mansoordaku/ckdisease

[2] H. G. M. Rautaray, M. K. Gourisaria, S. S. Rautaray and M. Pandey, "Pneumonia Detection Using CNN Through Chest X-Ray," *Journal of Engineering Science and Technology*, vol. 16, no. 1, pp. 861–876, Feb. 2021. [Online]. Available: https://jestec.taylors.edu.my/Vol

[3] P. Rajpurkar et al., "CheXNet: Radiologist-Level Pneumonia Detection on Chest X-Rays with Deep Learning," arXiv:1711.05225, 2017.

[4] D. Kermany et al., "Identifying Medical Diagnoses and Treatable Diseases by Image-Based Deep Learning," *Cell*, vol. 172, no. 5, pp. 1122–1131, 2018.

[5] G. Litjens et al., "A Survey on Deep Learning in Medical Image Analysis," *Medical Image Analysis*, vol. 42, pp. 60–88, 2017.

[6] S. Tajbakhsh et al., "Convolutional Neural Networks for Medical Image Analysis: Full Training or Fine Tuning?" *IEEE Transactions on Medical Imaging*, vol. 35, no. 5, pp. 1299–1312, 2016.

[7] K. He, X. Zhang, S. Ren and J. Sun, "Deep Residual Learning for Image Recognition," *IEEE CVPR*, 2016.

[8] G. Huang, Z. Liu, L. Van Der Maaten and K. Q. Weinberger, "Densely Connected Convolutional Networks," *IEEE CVPR*, 2017.

[9] M. Tan and Q. Le, "EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks," *ICML*, 2019.

[10] J. Hu, L. Shen and G. Sun, "Squeeze-and-Excitation Networks," *IEEE CVPR*, 2018.

[11] A. Howard et al., "MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications," arXiv:1704.04861, 2017.

[12] R. R. Selvaraju et al., "Grad-CAM: Visual Explanations from Deep Networks via Gradient-Based Localization," *IEEE ICCV*, 2017.

[13] L. Wang and A. Wong, "COVID-Net: A Tailored Deep Convolutional Neural Network Design for Detection of COVID-19 Cases from Chest X-Ray Images," *Scientific Reports*, 2020.

[14] F. Chollet, "Xception: Deep Learning with Depthwise Separable Convolutions," *IEEE CVPR*, 2017.

[15] S. Ioffe and C. Szegedy, "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift," *ICML*, 2015.

[16] N. Srivastava et al., “Dropout: A Simple Way to Prevent Neural Networks from Overfitting,” *Journal of Machine Learning Research*, vol. 15, pp. 1929–1958, 2014.
[17] O. Oktay et al., “Attention U-Net: Learning Where to Look for the Pancreas,” *MICCAI*, 2018.
[18] M. Tan and Q. Le, “EfficientNetV2: Smaller Models and Faster Training,” *ICML*, 2021.